

# **Runtime Environments**

CS143

Lecture 11

Instructor: Fredrik Kjolstad

Slide design by Prof. Alex Aiken, with modifications

# Status

---

- We have covered the front-end phases
  - Lexical analysis
  - Parsing
  - Semantic analysis
- Next are the back-end phases
  - Optimization
  - Code generation
- We'll do code generation first . . .

# Run-time environments

---

- Before discussing code generation, we need to understand what we are trying to generate
- There are a number of standard techniques for structuring executable code that are widely used

# Outline

---

- Management of run-time resources
- Correspondence between
  - static (compile-time) and
  - dynamic (run-time) structures
- Storage organization

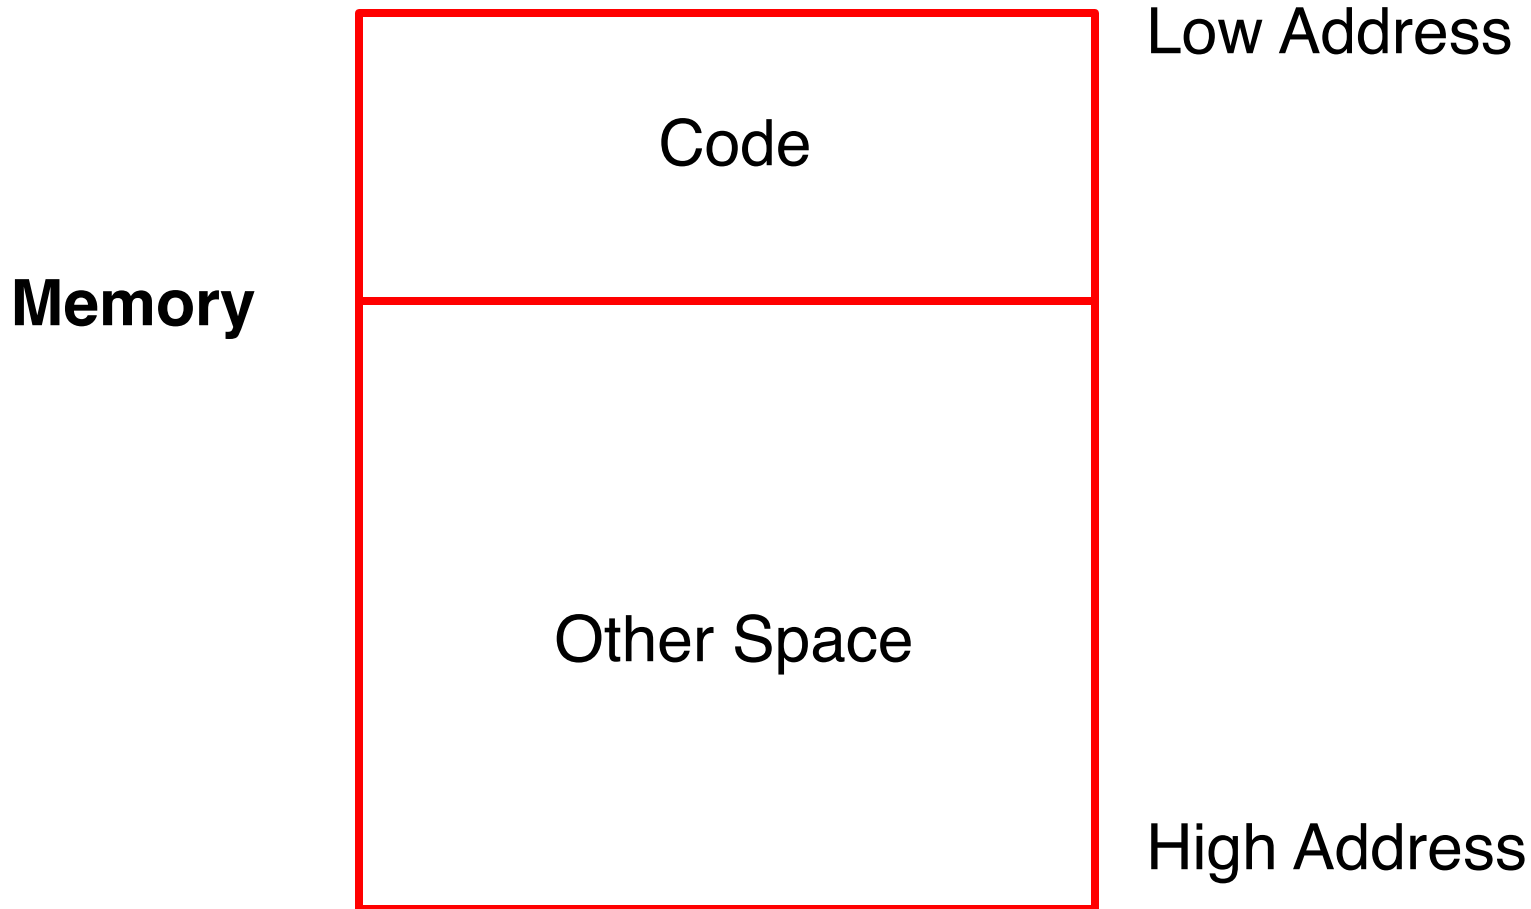
# Run-time Resources

---

- Execution of a program is initially under the control of the operating system
- When a program is invoked:
  - The OS allocates space for the program
  - The code is loaded into part of the space
  - The OS jumps to the entry point (i.e., “main”)

# Memory Layout

---



# Notes

---

- By tradition, pictures of machine organization have:
  - Low address at the top
  - High address at the bottom
  - Lines delimiting areas for different kinds of data
- These pictures are simplifications
  - E.g., not all memory need be contiguous

# What is Other Space?

---

- Holds all data for the program
- Other Space = Data Space
- Compiler is responsible for:
  - Generating code
  - Orchestrating use of the data area



# Code Generation Goals

---

- Two goals:
  - Correctness
  - Speed
- Most complications in code generation come from trying to be fast as well as correct

# Assumptions about Execution

---

1. Execution is sequential; control moves from one point in a program to another in a well-defined order
2. When a procedure is called, control eventually returns to the point immediately after the call

Do these assumptions always hold?

# Activations

---

- An invocation of procedure  $P$  is an activation of  $P$
- The lifetime of an activation of  $P$  is
  - All the steps to execute  $P$
  - Including all the steps in procedures  $P$  calls

# Lifetimes of Variables

---

- The lifetime of a variable **x** is the portion of execution in which **x** is defined
- Note that
  - Lifetime is a dynamic (run-time) concept
  - Scope is a static concept

# Activation Trees

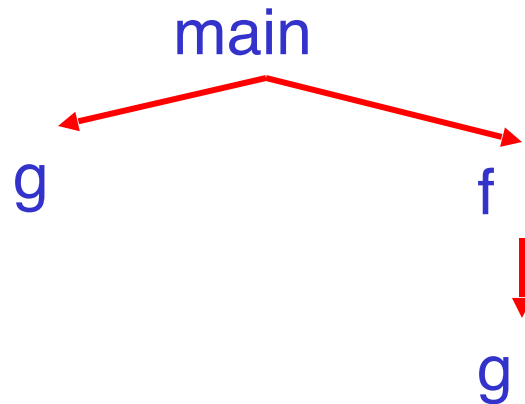
---

- Assumption (2) requires that when **P** calls **Q**, then **Q** returns before **P** does
- Lifetimes of procedure activations are properly nested
- Activation lifetimes can be depicted as a tree

# Example

---

```
Class Main {  
  g() : Int { 1 };  
  f(): Int { g() };  
  main(): Int {{ g(); f(); }};  
}
```



## Example 2

---

```
Class Main {  
    g() : Int { 1 };  
    f(x:Int): Int { if x = 0 then g() else f(x - 1) fi };  
    main(): Int { f(3) };  
}
```

What is the activation tree for this example?

# Notes

---

- The activation tree depends on run-time behavior
- The activation tree may be different for every program input
- Since activations are properly nested, a stack can track currently active procedures



# Example

---

```
Class Main {  
  g() : Int { 1 };  
  f(): Int { g() };  
  main(): Int {{ g(); f(); }};  
}
```

main

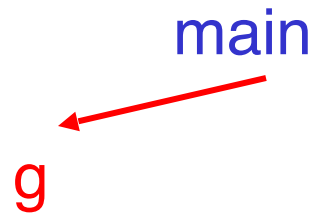
**Stack**

main

# Example

---

```
Class Main {  
  g() : Int { 1 };  
  f(): Int { g() };  
  main(): Int {{ g(); f(); }};  
}
```



**Stack**

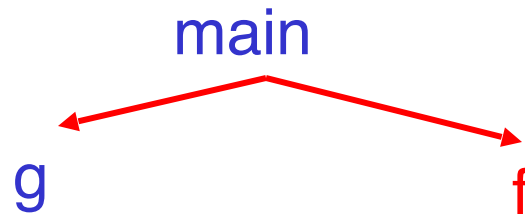
main

g

# Example

---

```
Class Main {  
  g() : Int { 1 };  
  f(): Int { g() };  
  main(): Int {{ g(); f(); }};  
}
```



**Stack**

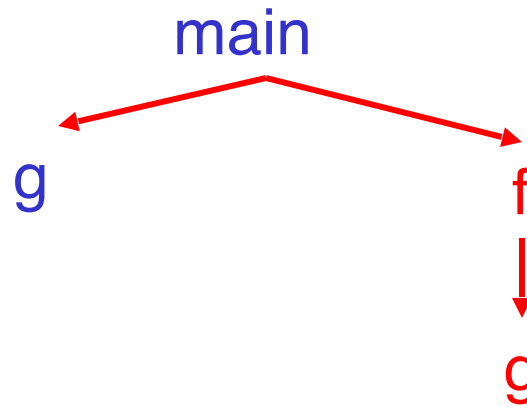
main

f

# Example

---

```
Class Main {  
  g() : Int { 1 };  
  f(): Int { g() };  
  main(): Int {{ g(); f(); }};  
}
```



## Stack

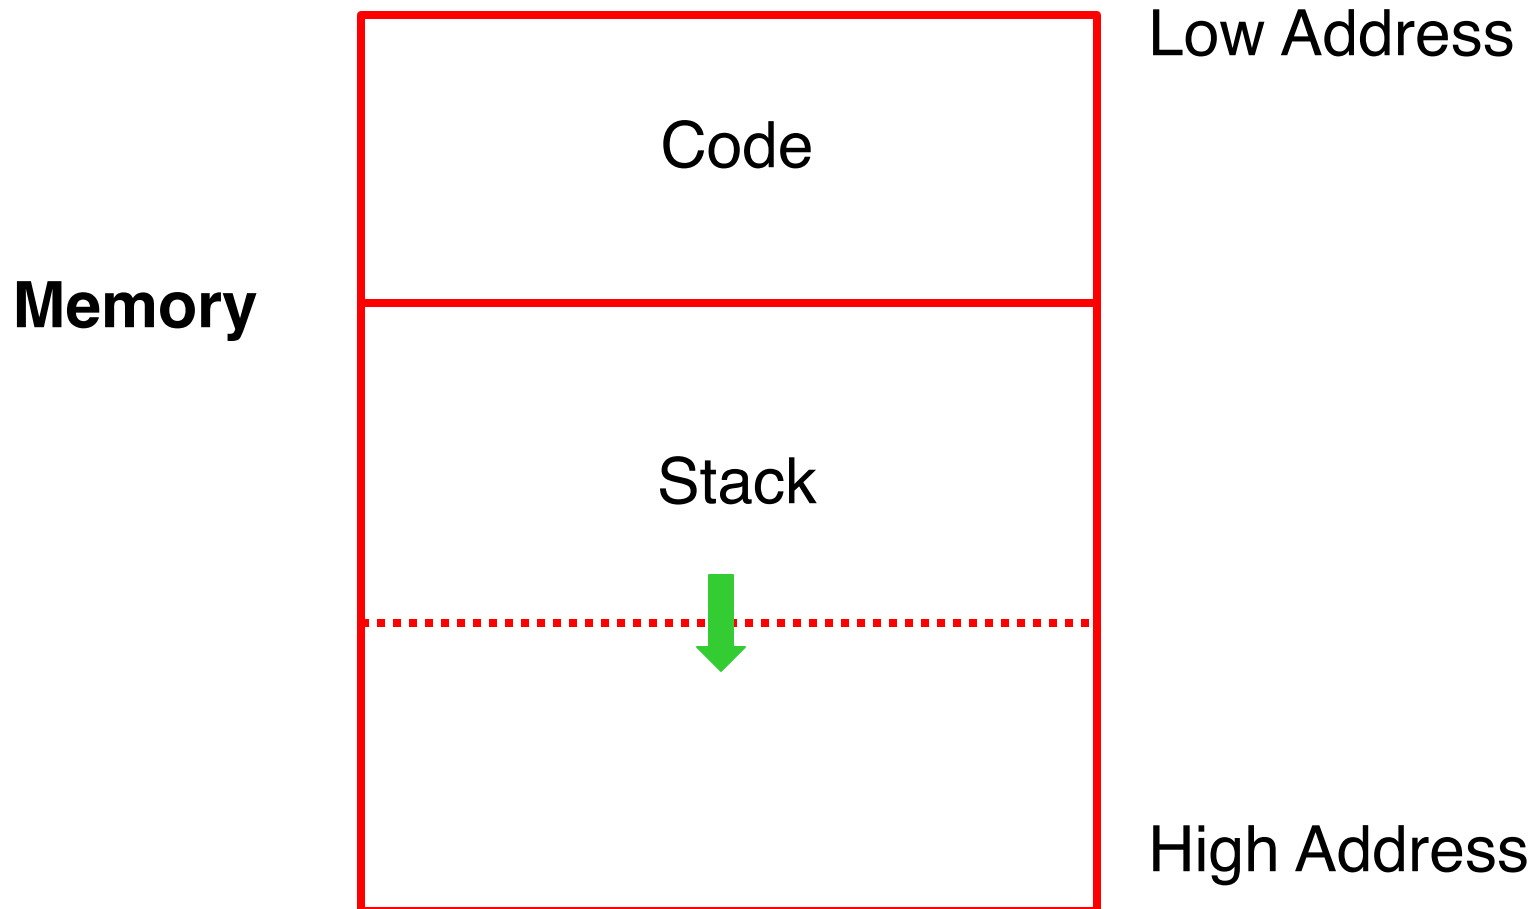
main

f

g

# Revised Memory Layout

---



# Activation Records

---

- The information needed to manage one procedure activation is called an activation record (AR) or a stack frame
- If procedure **F** calls **G**, then **G**'s activation record contains a mix of info about **F** and **G**.

## What is in **G**'s AR when **F** calls **G**?

---

- **F** is “suspended” until **G** completes, at which point **F** resumes. **G**'s AR contains information needed to resume execution of **F**.
- **G**'s AR may also contain:
  - **G**'s return value (needed by **F**)
  - Actual parameters to **G** (supplied by **F**)
  - Space for **G**'s local variables

# The Contents of a Typical AR for G

---

- Space for G's return value
- Actual parameters
- Pointer to the previous activation record
  - The control link; points to AR of caller of G
- Machine status prior to calling G
  - Contents of registers & program counter
  - Local variables
- Other temporary values



## Example 2, Revisited

---

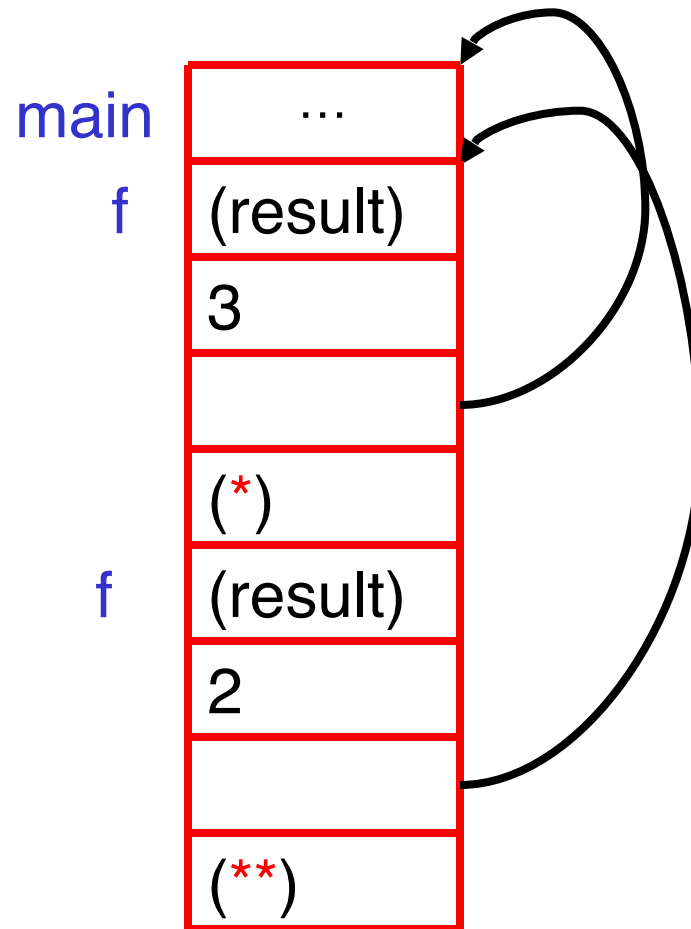
```
Class Main {  
  g() : Int { 1 };  
  f(x:Int):Int { if x=0 then g() else f(x - 1)(**) fi };  
  main(): Int { f(3)(*) };  
}
```

AR for f:

result
argument
control link
return address

# Stack After Two Calls to **f**

---



# Notes

---

- `main` has no argument or local variables and its result is never used; its AR is uninteresting
- `(*)` and `(**)` are return addresses of the invocations of `f`
  - The return address is where execution resumes after a procedure call finishes
- This is only one of many possible AR designs
  - Would also work for C, Pascal, FORTRAN, etc.

# The Main Point

---

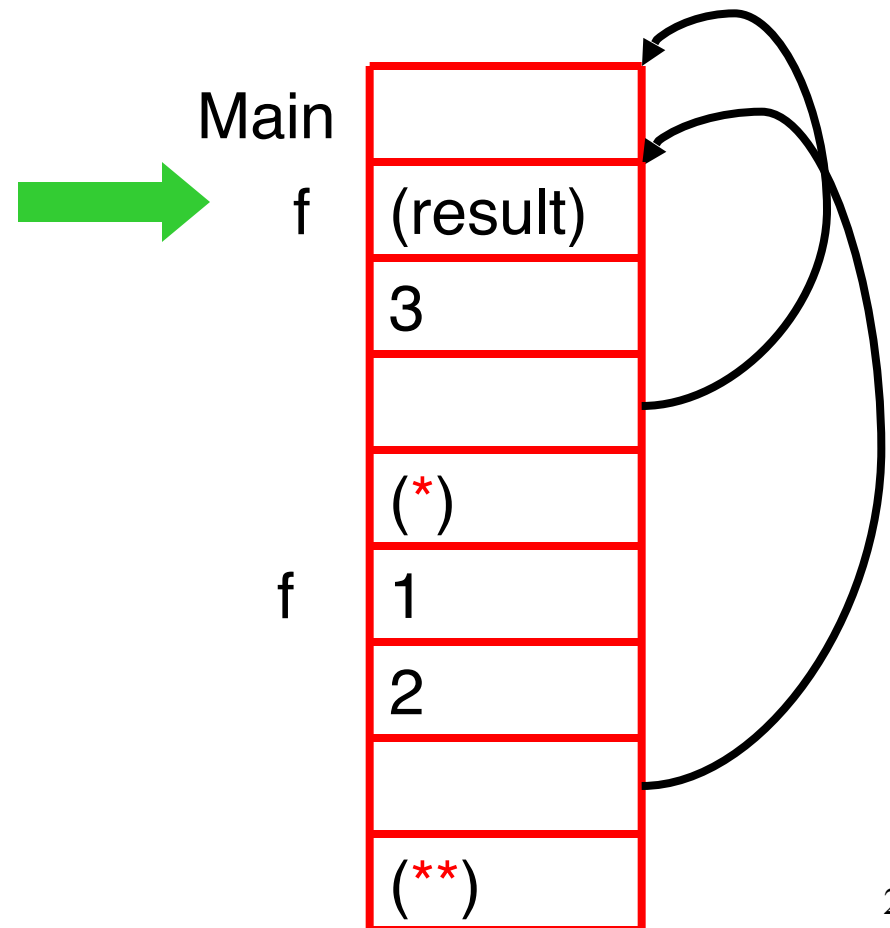
The compiler must determine, at compile-time, the layout of activation records and generate code that correctly accesses locations in the activation record

Thus, the AR layout and the code generator must be designed together!

# Example

---

The picture shows the state after the call to the 2nd invocation of **f** returns



# Discussion

---

- The advantage of placing the return value 1st in a frame is that the caller can find it at a fixed offset from its own frame
- There is nothing magic about this organization
  - Can rearrange order of frame elements
  - Can divide caller/callee responsibilities differently
  - An organization is better if it improves execution speed or simplifies code generation

## Discussion (Cont.)

---

- Real compilers hold as much of the frame as possible in registers
  - Especially the method result and arguments

# Globals

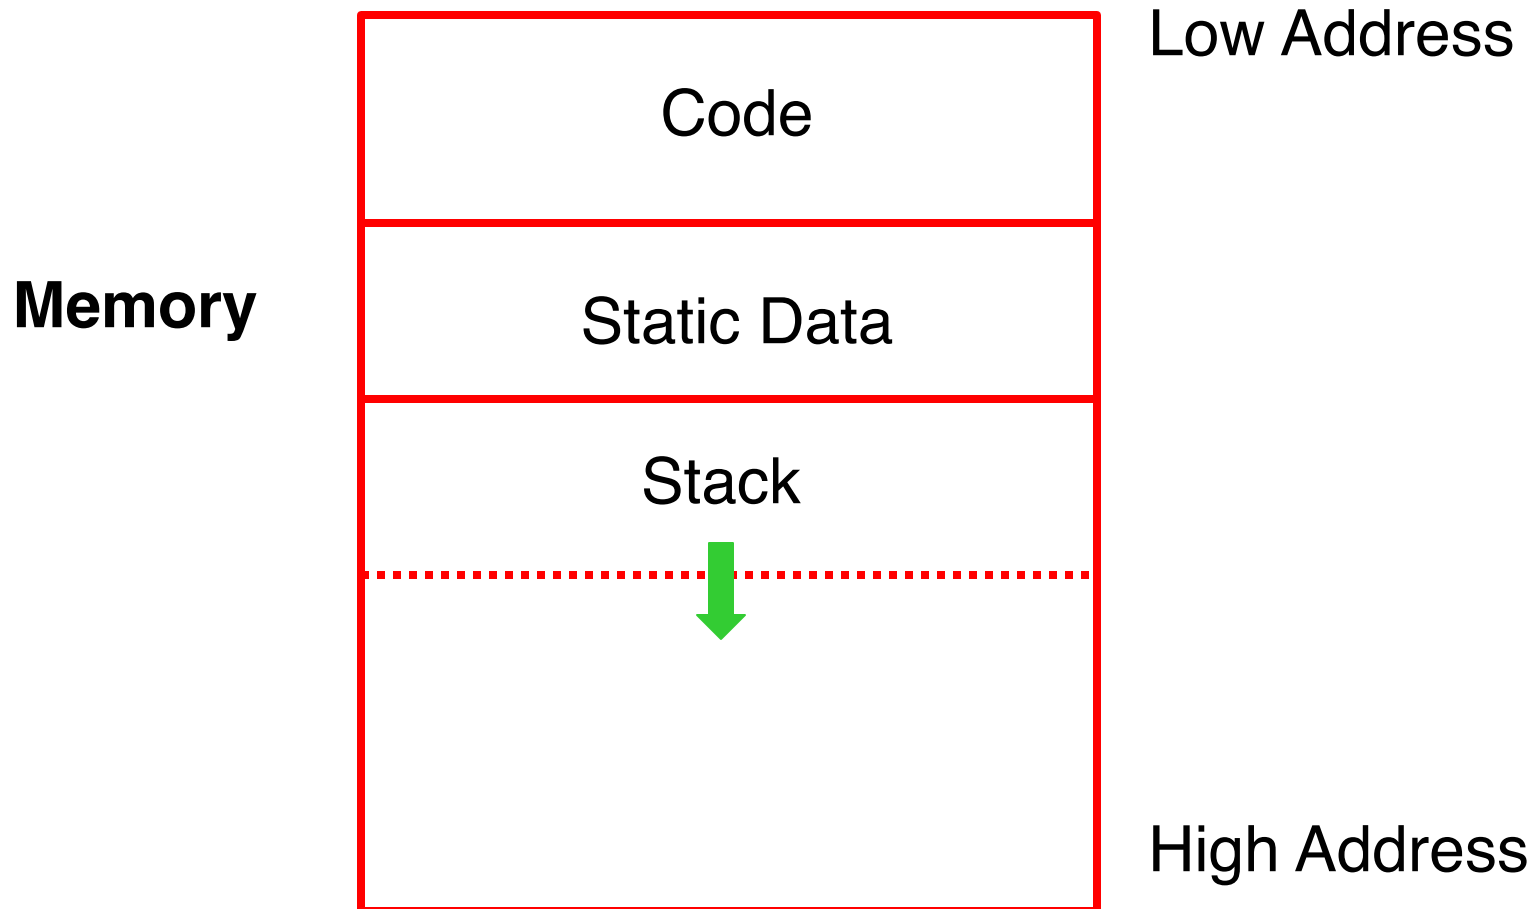
---

- All references to a global variable point to the same object
  - Can't store a global in an activation record
- Globals are assigned a fixed address once
  - Variables with fixed address are “statically allocated”
- Depending on the language, there may be other statically allocated values



# Memory Layout with Static Data

---



# Heap Storage

---

- A value that outlives the procedure that creates it cannot be kept in the AR

`method foo() { new Bar }`

The `Bar` value must survive deallocation of `foo`'s AR

- Languages with dynamically allocated data use a heap to store dynamic data

# Notes

---

- The code area contains object code
  - For most languages, fixed size and read only
- The static area contains data (not code) with fixed addresses (e.g., global data)
  - Fixed size, may be readable or writable
- The stack contains an AR for each currently active procedure
  - Each AR usually fixed size, contains locals
- Heap contains all other data
  - In C, heap is managed by malloc and free

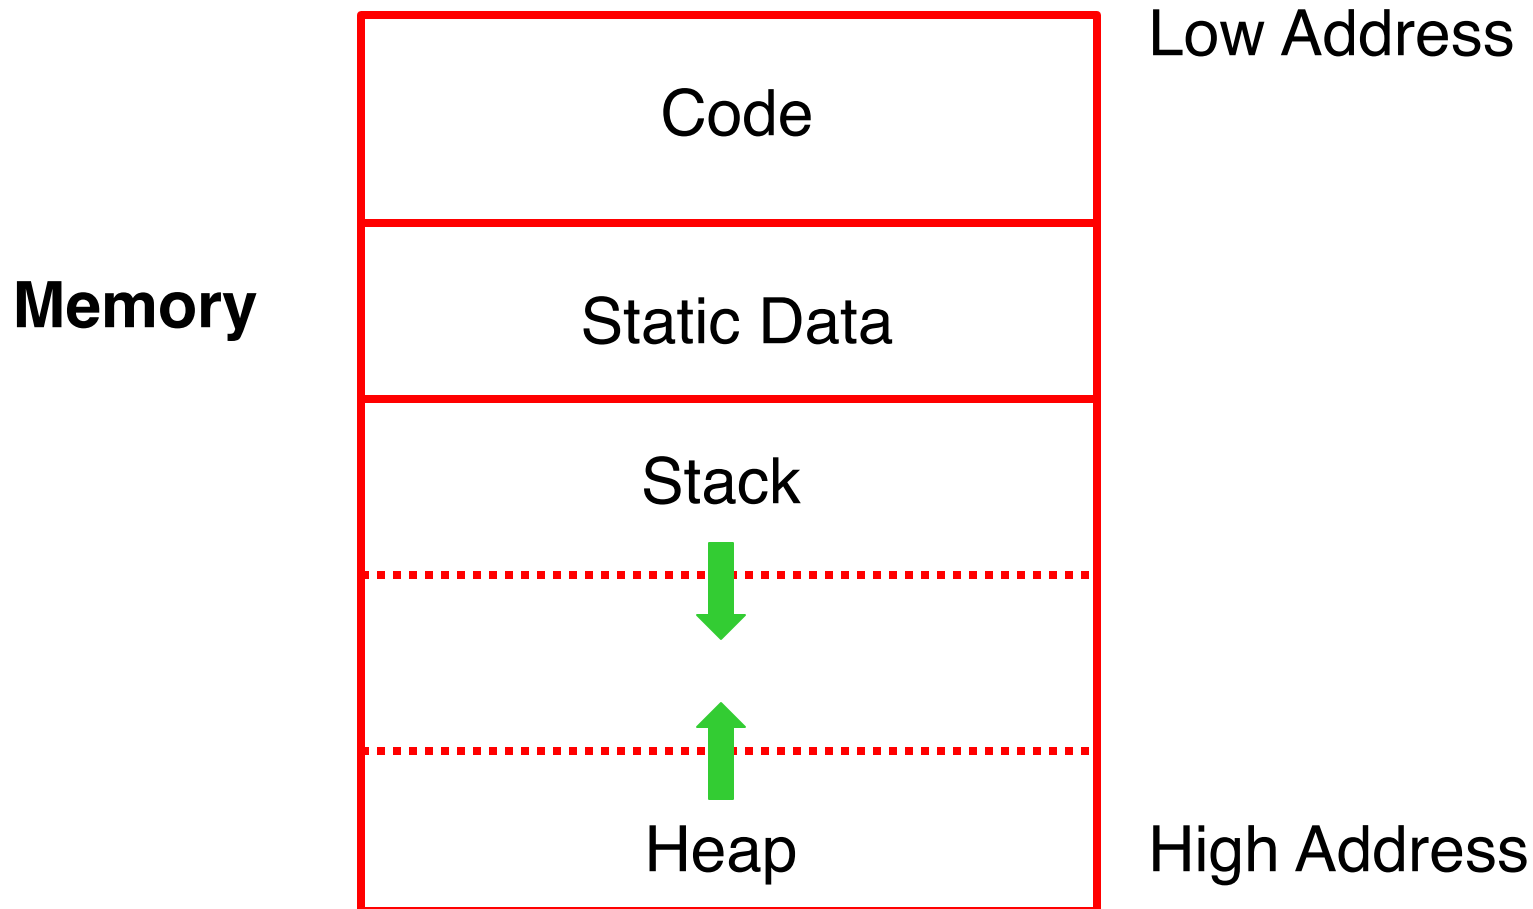
## Notes (Cont.)

---

- Both the heap and the stack grow
- Must take care that they don't grow into each other
- Solution: start heap and stack at opposite ends of memory and let them grow towards each other

# Memory Layout with Heap

---



# Data Layout

---

- Low-level details of machine architecture are important in laying out data for correct code and maximum performance
- Chief among these concerns is alignment

# Alignment

---

- Most machines are 32 or 64 bit
  - 8 bits in a byte
  - 4 bytes in a word
  - Machines are either byte or word addressable
- Data is word aligned if it begins at a word boundary
- Most machines have some alignment restrictions
  - Or performance penalties for poor alignment

## Alignment (Cont.)

---

- Example: A string  
“Hello”  
Takes 5 characters (without a terminating \0)
- To word align next datum, add 3 “padding” characters to the string
- The padding is not part of the string, it’s just unused memory



## Next Topic: Stack Machines

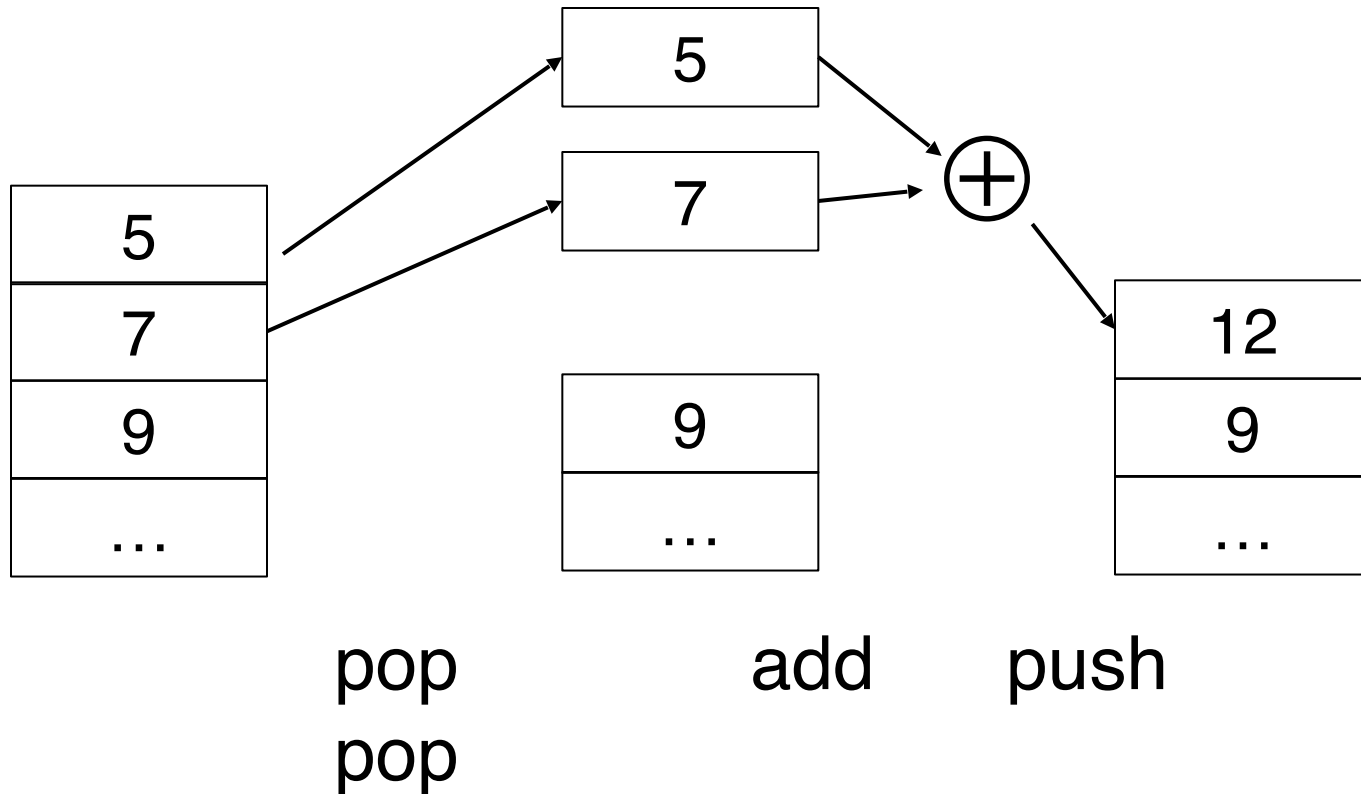
---

- A simple evaluation model for expressions
- No variables or registers
- A stack of values for intermediate results
- Each instruction:
  - Takes its operands from the top of the stack
  - Removes those operands from the stack
  - Computes the required operation on them
  - Pushes the result on the stack

# Example of Stack Machine Operation

---

- The addition operation on a stack machine



# Example of a Stack Machine Program

---

- Consider two instructions
  - `push i` - place the integer `i` on top of the stack
  - `add` - pop two elements, add them and put the result back on the stack
- A program to compute  $7 + 5$ :
  - `push 7`
  - `push 5`
  - `add`

# Why Use a Stack Machine?

---

- Each operation takes operands from the same place and puts results in the same place
- This means a uniform compilation scheme
- And therefore a simpler compiler

# Why Use a Stack Machine?

---

- Location of the operands is implicit
  - Always on the top of the stack
- No need to specify operands explicitly
- No need to specify the location of the result
- Instruction “**add**” as opposed to “**add**  $r_1$ ,  $r_2$ ”
  - ⇒ Smaller encoding of instructions
  - ⇒ More compact programs
- This is one reason why Java Bytecodes and WebAssembly use a stack evaluation model

# Optimizing the Stack Machine

---

- The add instruction does 3 memory operations
  - Two reads and one write to the stack
  - The top of the stack is frequently accessed
- Idea: keep the top of the stack in a register (called accumulator)
  - Register accesses are faster
- The “add” instruction is now
$$\text{acc} \leftarrow \text{acc} + \text{top\_of\_stack}$$
  - Only one memory operation!

# Stack Machine with Accumulator

---

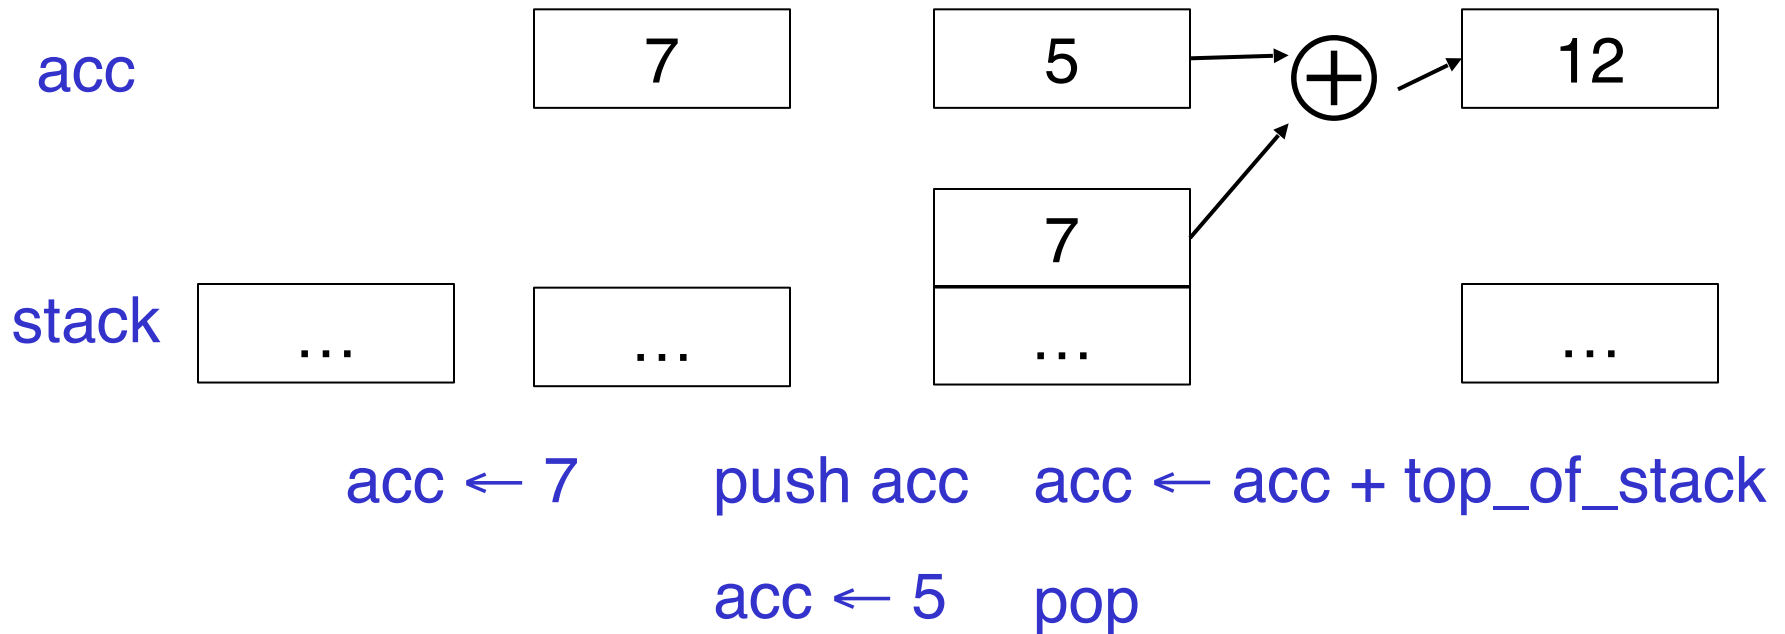
## Invariants

- The result of an expression is in the accumulator
- For  $op(e_1, \dots, e_n)$ 
  - Evaluate each of  $e_1, \dots, e_{n-1}$  and push results on the stack
  - Evaluating  $e_n$  leaves the result in the accumulator
  - Evaluating  $op$  pops  $n-1$  values from the stack and updates the accumulator
- Expression evaluation preserves the stack

# Stack Machine with Accumulator. Example

---

- Compute  $7 + 5$  using an accumulator





## A Bigger Example: $3 + (7 + 5)$

---

Code	Acc	Stack
$\text{acc} \leftarrow 3$	3	<init>
push acc	3	3, <init>
$\text{acc} \leftarrow 7$	7	3, <init>
push acc	7	7, 3, <init>
$\text{acc} \leftarrow 5$	5	7, 3, <init>
$\text{acc} \leftarrow \text{acc} + \text{top\_of\_stack}$	12	7, 3, <init>
pop	12	3, <init>
$\text{acc} \leftarrow \text{acc} + \text{top\_of\_stack}$	15	3, <init>
pop	15	<init>

# Notes

---

- It is very important evaluation of a subexpression preserves the stack
  - Stack before the evaluation of  $7 + 5$  is  $3, \langle \text{init} \rangle$
  - Stack after the evaluation of  $7 + 5$  is  $3, \langle \text{init} \rangle$
  - The first operand is on top of the stack